
Cloud First Consulting

Cloud and AI Platforms

AI application platforms

A successful AI prototype can still lack the controls needed for a dependable service. Establish a deployable platform for models, knowledge, and tools with explicit evaluation, release, monitoring, and cost responsibilities before production use.

Capabilities

- Select hosting and supported model options against quality, latency, region, and cost requirements.
- Design identity, secrets handling, knowledge access, and approved Model Context Protocol integrations.
- Define orchestration boundaries, workflow handoffs, retries, and human escalation paths.
- Create representative evaluations and release gates for quality, access controls, and failure behavior.
- Use GitHub Copilot for development assistance with human code review and tests.
- Plan budgets, monitoring, recovery, version changes, and accountable support.

Deliverables

- A reference architecture with hosting choices, trust boundaries, and documented tradeoffs.
- A scoped platform foundation with environment configuration and repeatable deployment steps.
- An evaluation and release plan with test datasets, acceptance thresholds, and rollback criteria.
- An operations runbook covering monitoring, cost alerts, incidents, and model or application changes.

Business value

Make the transition from prototype to operation a deliberate engineering decision. Baseline task quality, security findings, latency, availability objectives, and cost per useful task; release only when agreed readiness criteria are met.

Start with a use case and a production-readiness review.

Bring a representative application, data and model requirements, intended integrations, and operating constraints. Agree the platform boundary and pilot workload; production readiness remains separate from prototype performance.

Foundry, Azure, Entra, GitHub, GitHub Copilot

© 2026 Cloud First Consulting